AI Report

We are moderately confident this text is

AI Generated

AI Probability

100%

This number is the probability that the document is AI generated, not a percentage of AI text in the document.

Plagiarism

?

The plagiarism scan was not run for this document. Go to gptzero.me to check for plagiarism.

This paper points out that existing generative recommendation research primarily focuses on the sema - 2/16/2026

This paper points out that existing generative recommendation research primarily focuses on the semantic encoding of items and aligning item concepts with large language model representations. It proposes a CoFiGR framework, which includes bidirectional relational semantic collaborative IDs, multi-task continual pre-training, and cascade-prefix policy optimization. This framework aligns item-level, user-level, and interaction-level objectives through a coarse-to-fine process and fuses structured identifiers with natural language semantics. Experiments on two benchmark datasets and one industrial dataset validate the effectiveness of CoFiGR.

The research focused on a cutting-edge topic, i.e., generative recommendation systems. Combining various recommendation signals with large language models from a more comprehensive perspective is a promising approach.

Experimental evaluations were conducted on both publicly available datasets and industrial datasets.

The proposed solution offers only incremental technical novelty, aiming to make minor improvements to different steps of the generative recommendation process. For example, this involves introducing collaborative statistical information during the tokenization process, incorporating more objectives during the supervised fine-tuning phase, and implementing curriculum learning during the reinforcement learning phase.

The description of the proposed solution lacks sufficient detail and is difficult to understand. For example, 1) In Section 3.1, what are the high-frequency co-occurring item pairs? Are user identifiers represented solely by strings from the actions? 2) In Section 3.2, the specific process for all optimization objectives is very unclear, especially how user identifiers are used. 3) In Section 3.3, the specific process of the entire operation is also very unclear.

The experimental setup and results need to be improved. Some key issues are: 1) the lack of experimental analysis on time complexity to demonstrate whether the gains of CoFiGR are cost-effective; 2) a comparison of parameter sizes between CoFiGR and all baselines to determine whether the gains of CoFiGR stem from the advantages

of the model architecture itself; and 3) providing case studies to highlight the impact of certain design choices, such as the changes introduced by incorporating more information during the tokenization phase.

The detailed description of the solution needs to be revised to improve readability.

The reported metrics in the alignment experiments are inconsistent. Tables 3 through 5 report the results for Precision (unknown K) and Recall@20. This is inconsistent with the metric choices in Table 2. This is very confusing.

FAQs

What is GPTZero?

GPTZero is the leading AI detector for checking whether a document was written by a large language model such as ChatGPT. GPTZero detects AI on sentence, paragraph, and document level. Our model was trained on a large, diverse corpus of human-written and AI-generated text with support for English, Spanish, French, German, and other languages. To date, GPTZero has served over 10 million users around the world, and works with over 100 organizations in education, hiring, publishing, legal, and more.

When should I use GPTZero?

Our users have seen the use of AI-generated text proliferate into education, certification, hiring and recruitment, social writing platforms, disinformation, and beyond. We've created GPTZero as a tool to highlight the possible use of AI in writing text. In particular, we focus on classifying AI use in prose. Overall, our classifier is intended to be used to flag situations in which a conversation can be started (for example, between educators and students) to drive further inquiry and spread awareness of the risks of using AI in written work.

Does GPTZero only detect ChatGPT outputs?

No, GPTZero works robustly across a range of AI language models, including but not limited to ChatGPT, GPT-5, GPT-4, GPT-3, Gemini, Claude, and AI services based on those models.

What are the limitations of the classifier?

The nature of AI-generated content is changing constantly. As such, these results should not be used to punish students. We recommend educators to use our behind-the-scene Writing Reports as part of a holistic assessment of student work. There always exist edge cases with both instances where AI is classified as human, and human is classified as AI. Instead, we recommend educators take approaches that give students the opportunity to demonstrate their understanding in a controlled environment and craft assignments that cannot be solved with AI. Our classifier is not trained to identify AI-generated text after it has been heavily modified after generation (although we estimate this is a minority of the uses for AI-generation at the moment). Currently, our classifier can sometimes flag other machine-generated or highly procedural text as AI-generated, and as such, should be used on more descriptive portions of text.

I'm an educator who has found AI-generated text by my students. What do I do?

Firstly, at GPTZero, we don't believe that any AI detector is perfect. There always exist edge cases with both instances where AI is classified as human, and human is classified as AI. Nonetheless, we recommend that educators can do the following when they get a positive detection: Ask students to demonstrate their understanding in a controlled environment, whether that is through an in-person assessment, or through an editor that can track their edit history (for instance, using our Writing Reports through Google Docs). Check out our list of several recommendations on types of assignments that are difficult to solve with AI.

Ask the student if they can produce artifacts of their writing process, whether it is drafts, revision histories, or brainstorming notes. For example, if the editor they used to write the text has an edit history (such as Google Docs), and it was typed out with several edits over a reasonable period of time, it is likely the student work is authentic. You can use GPTZero's Writing Reports to replay the student's writing process, and view signals that indicate the authenticity of the work.

See if there is a history of AI-generated text in the student's work. We recommend looking for a long-term pattern of AI use, as opposed to a single instance, in order to determine whether the student is using AI.